



Tu non puoi passare!

Policy compliance con OPA Gatekeeper



Niccolò Raspa - *Kubernetes Engineer @ SIGHUP*

KCD Italy  Nov 17-18, 2021



👋 Ciao

My name is Niccolò and I work at
SIGHUP as a Kubernetes Engineer.



linkedin.com/in/niccolo-raspa/



[@niccoloraspa](https://twitter.com/niccoloraspa)

Agenda





The Balrog and policy compliance



The bridge of Khazad-dûm and the road to ETCD




Gandalf and Opa Gatekeeper



The fall and the demo

What is a Balrog?

```
apiVersion: v1
kind: Pod
metadata:
  name: balrog
spec:
  securityContext:
    runAsUser: 0
    privileged: True
    capabilities:
      add: ["CAP_SYS_BOOT"]
  containers:
    - name: busybox
      image: evil.registry.io/balrog/busybox
      command: [ "sh", "-c", "sleep 1h" ]
      volumeMounts:
        - name: private-volume
          mountPath: /private
  volumes:
    - name: private-volume
      hostPath:
        path: /private
        type: Directory
```



kube api-server




etcd



The road to etcd



Kubernetes cluster



Authentication

Authorization

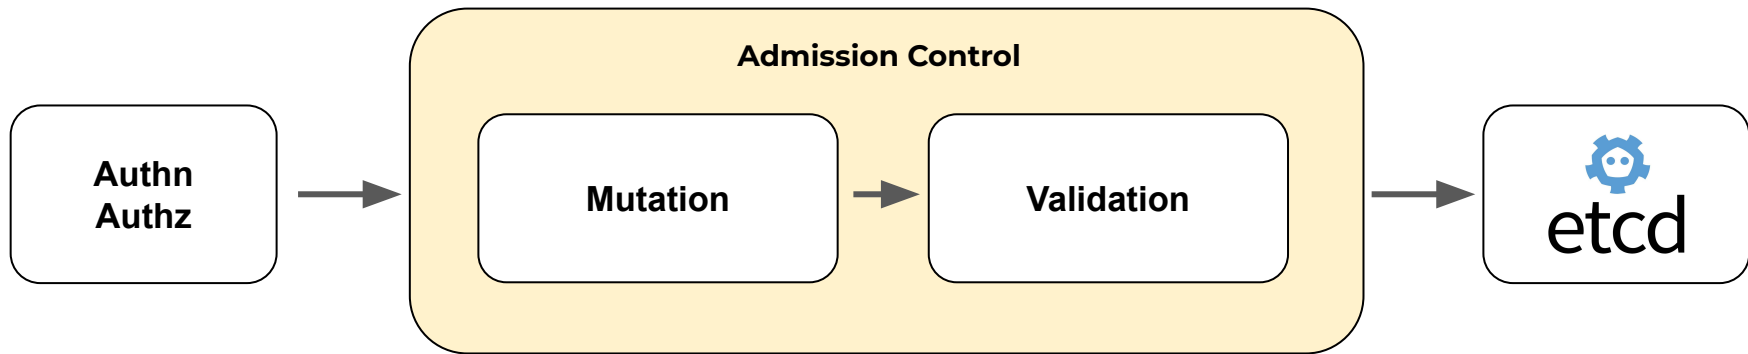
Admission control



Admission Control



An admission controller is a piece of code that processes requests to the Kubernetes API server prior to persistence of the object, but after the request is authenticated and authorized.



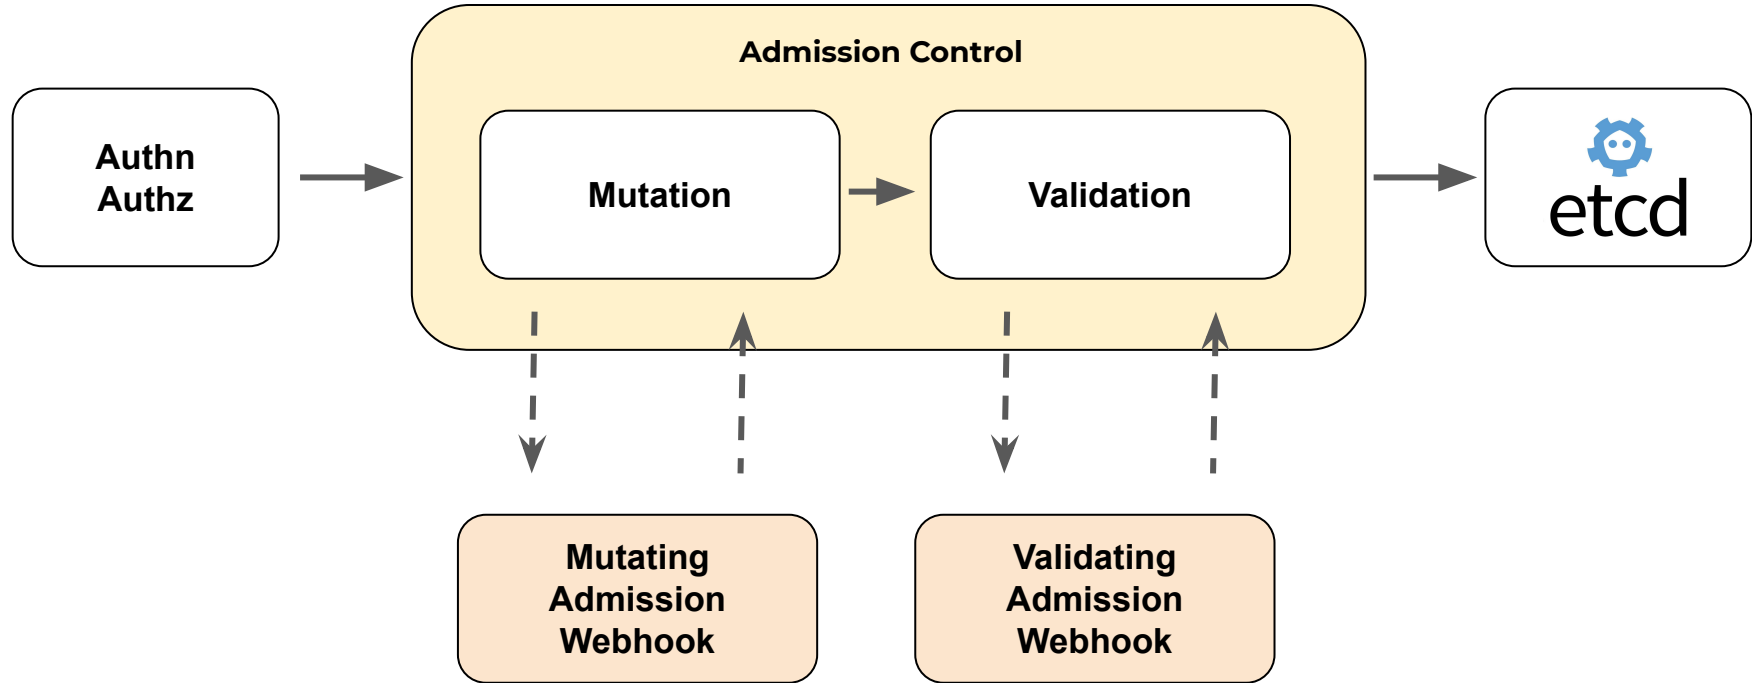
You are already using them:

```
kube-apiserver -h | grep enable-admission-plugins
```

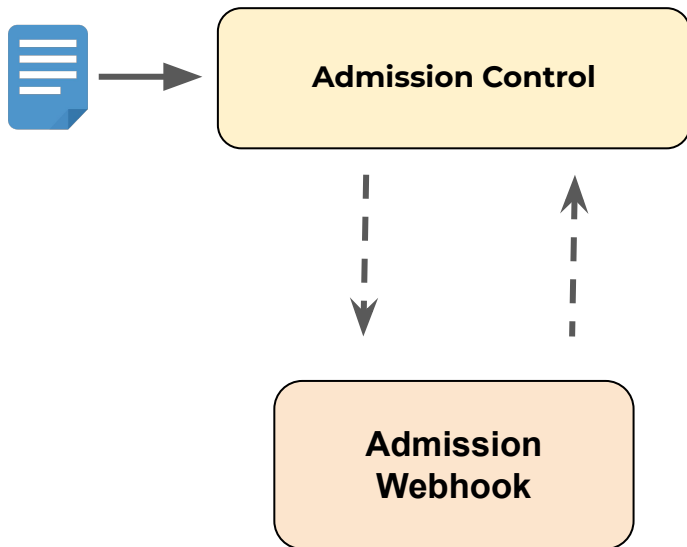
--enable-admission-plugins strings

admission plugins that should be enabled in addition to default enabled ones (NamespaceLifecycle, LimitRanger, ServiceAccount, TaintNodesByCondition, PodSecurity, Priority, DefaultTolerationSeconds, DefaultStorageClass, StorageObjectInUseProtection, PersistentVolumeClaimResize, RuntimeClass, CertificateApproval, CertificateSigning, CertificateSubjectRestriction, DefaultIngressClass, MutatingAdmissionWebhook, ValidatingAdmissionWebhook, ResourceQuota). Comma-delimited list of admission plugins: AlwaysAdmit, AlwaysDeny, AlwaysPullImages, CertificateApproval, CertificateSigning, CertificateSubjectRestriction, DefaultIngressClass, DefaultStorageClass, DefaultTolerationSeconds, DenyServiceExternalIPs, EventRateLimit, ExtendedResourceToleration, ImagePolicyWebhook, LimitPodHardAntiAffinityTopology, LimitRanger, MutatingAdmissionWebhook, NamespaceAutoProvision, NamespaceExists, NamespaceLifecycle, NodeRestriction, OwnerReferencesPermissionEnforcement, PersistentVolumeClaimResize, PersistentVolumeLabel, PodNodeSelector, PodSecurity, PodSecurityPolicy, PodTolerationRestriction, Priority, ResourceQuota, RuntimeClass, SecurityContextDeny, ServiceAccount, StorageObjectInUseProtection, TaintNodesByCondition, ValidatingAdmissionWebhook. The order of plugins in this flag does not matter.

Admission Webhooks

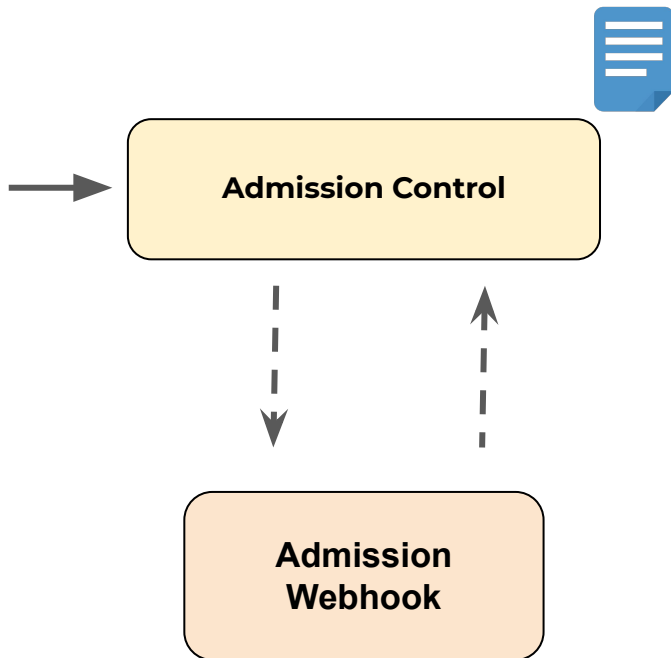


Admission Webhooks Flow



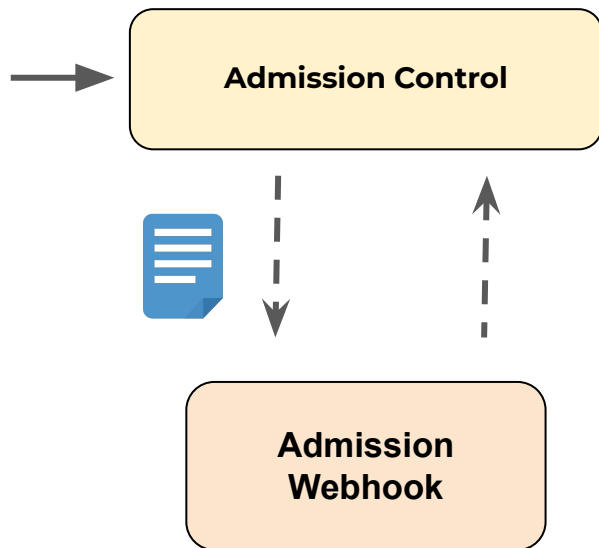
```
kind: Pod
apiVersion: v1
metadata:
  name: myapp
spec:
  containers:
    - image: nginx
      name: nginx-frontend
    - image: mysql
      name: mysql-backend
```

Admission Webhooks Flow



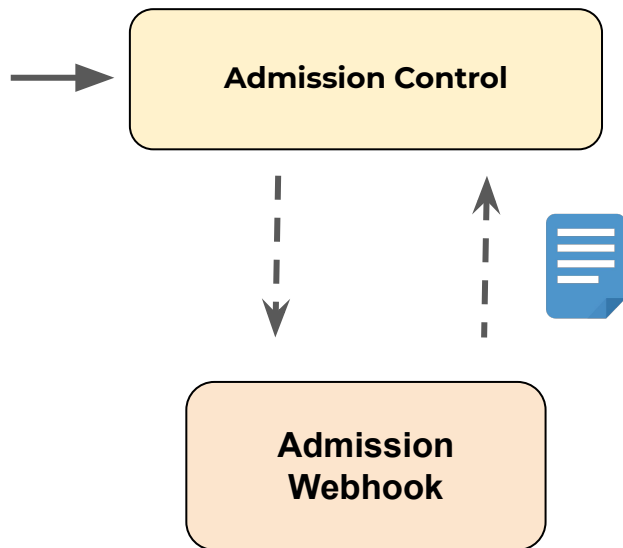
```
apiVersion: admissionregistration.k8s.io/v1
kind: ValidatingWebhookConfiguration
metadata:
  name: "pod-policy.example.com"
webhooks:
- name: "pod-policy.example.com"
  rules:
  - apiGroups:  [""]
    apiVersions: ["v1"]
    operations:  ["CREATE"]
    resources:   ["pods"]
    scope:       "Namespaced"
  clientConfig:
    service:
      namespace: "my-namespace"
      name: "my-admission-webhook-svc"
      caBundle: "Ci0tLS0tQk.."
  admissionReviewVersions: ["v1", "v1beta1"]
  sideEffects: None
  timeoutSeconds: 5
```

Admission Webhooks Flow



```
{
  "kind": "AdmissionReview",
  "request": {
    "kind": {
      "kind": "Pod",
      "version": "v1"
    },
    "object": {
      "metadata": {
        "name": "myapp"
      },
      "spec": {
        "containers": [
          {
            "image": "nginx",
            "name": "nginx-frontend"
          },
          {
            "image": "mysql",
            "name": "mysql-backend"
          }
        ]
      }
    }
  }
}
```

Admission Webhooks Flow

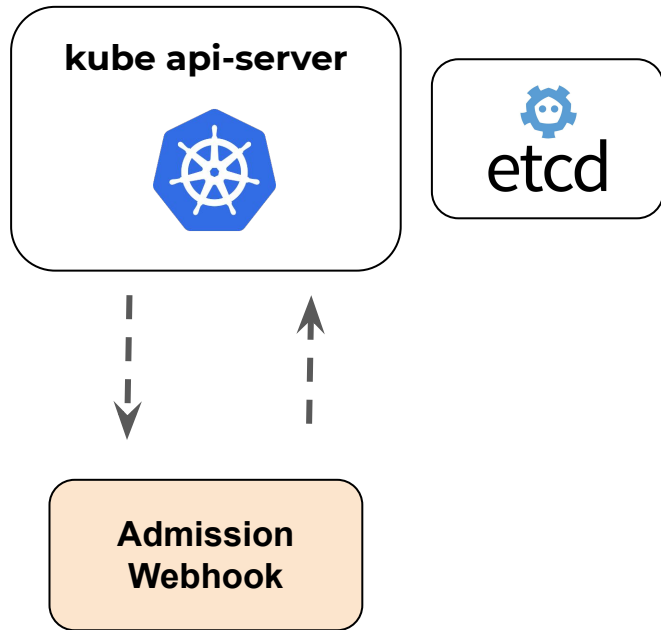



```
{
  "apiVersion": "admission.k8s.io/v1",
  "kind": "AdmissionReview",
  "response": {
    "uid": "<value from request.uid>",
    "allowed": true
  }
}
```

What's wrong with this approach?

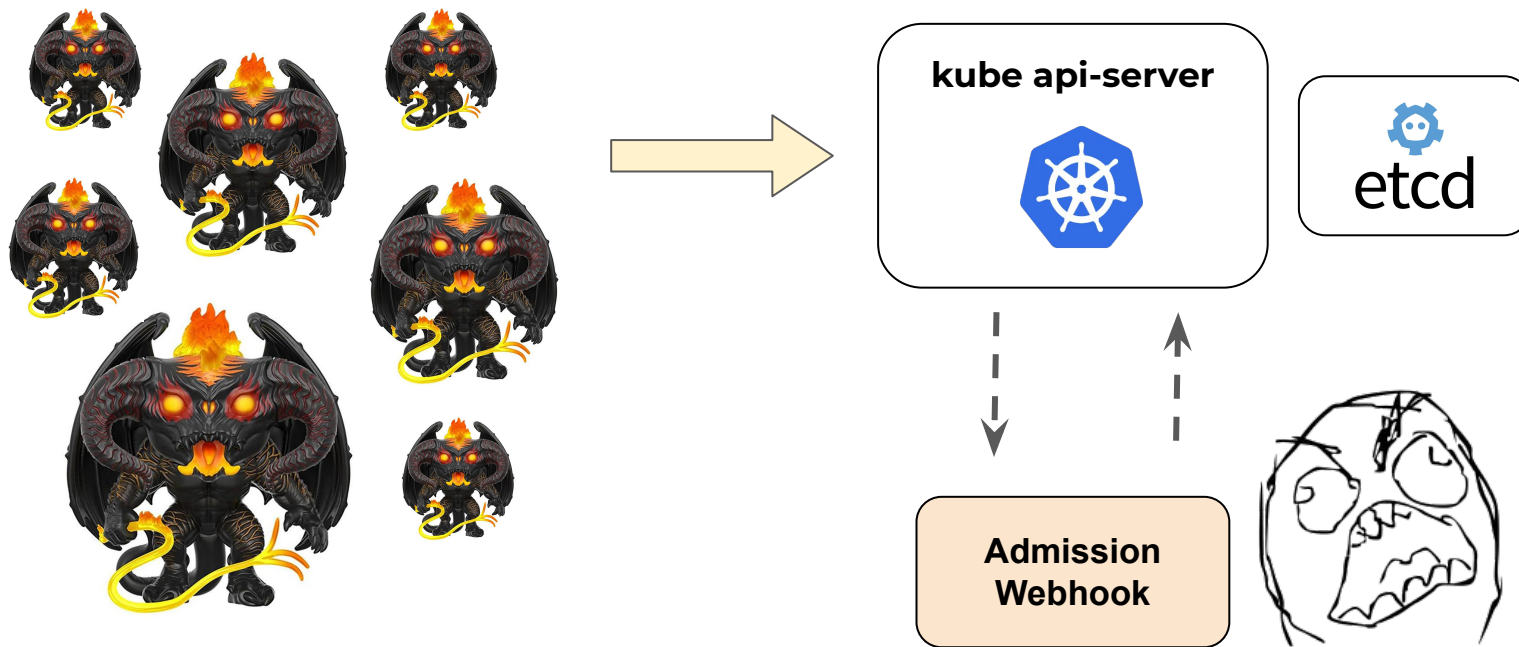


Kubernetes allows **decoupling policy decisions** from the API server with admission controller webhooks intercepting admission requests before they are persisted.

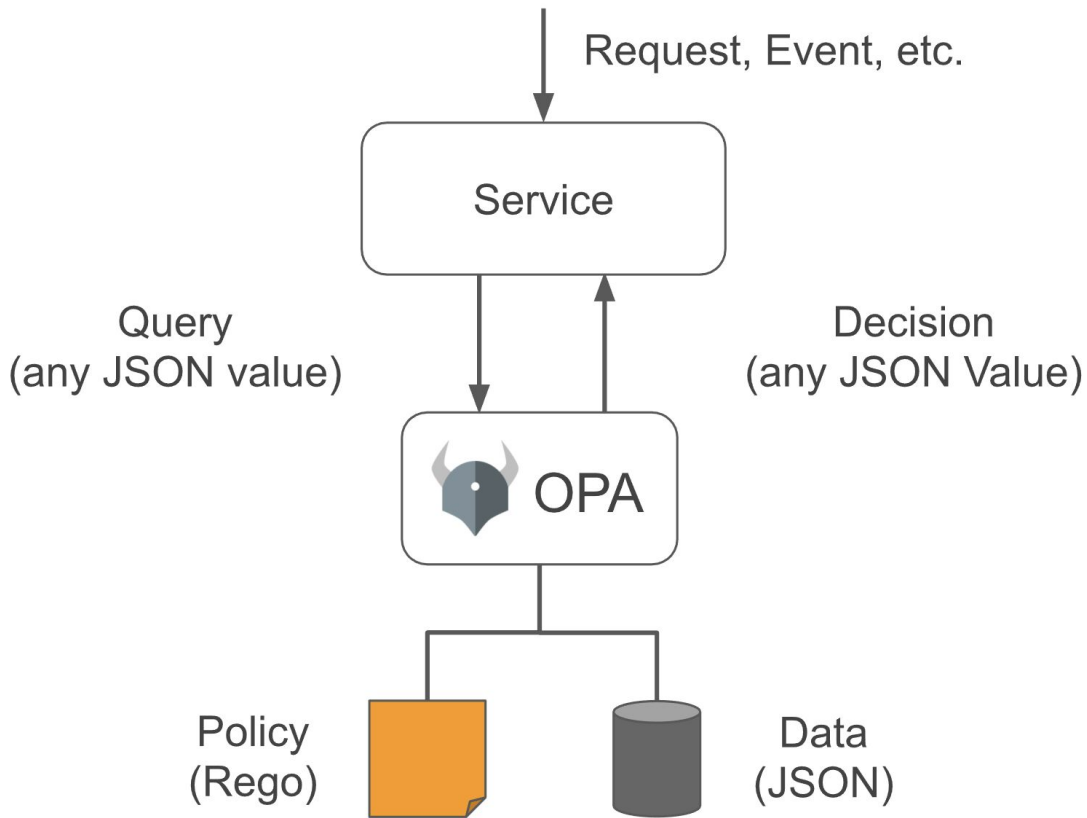


What's wrong with this approach?

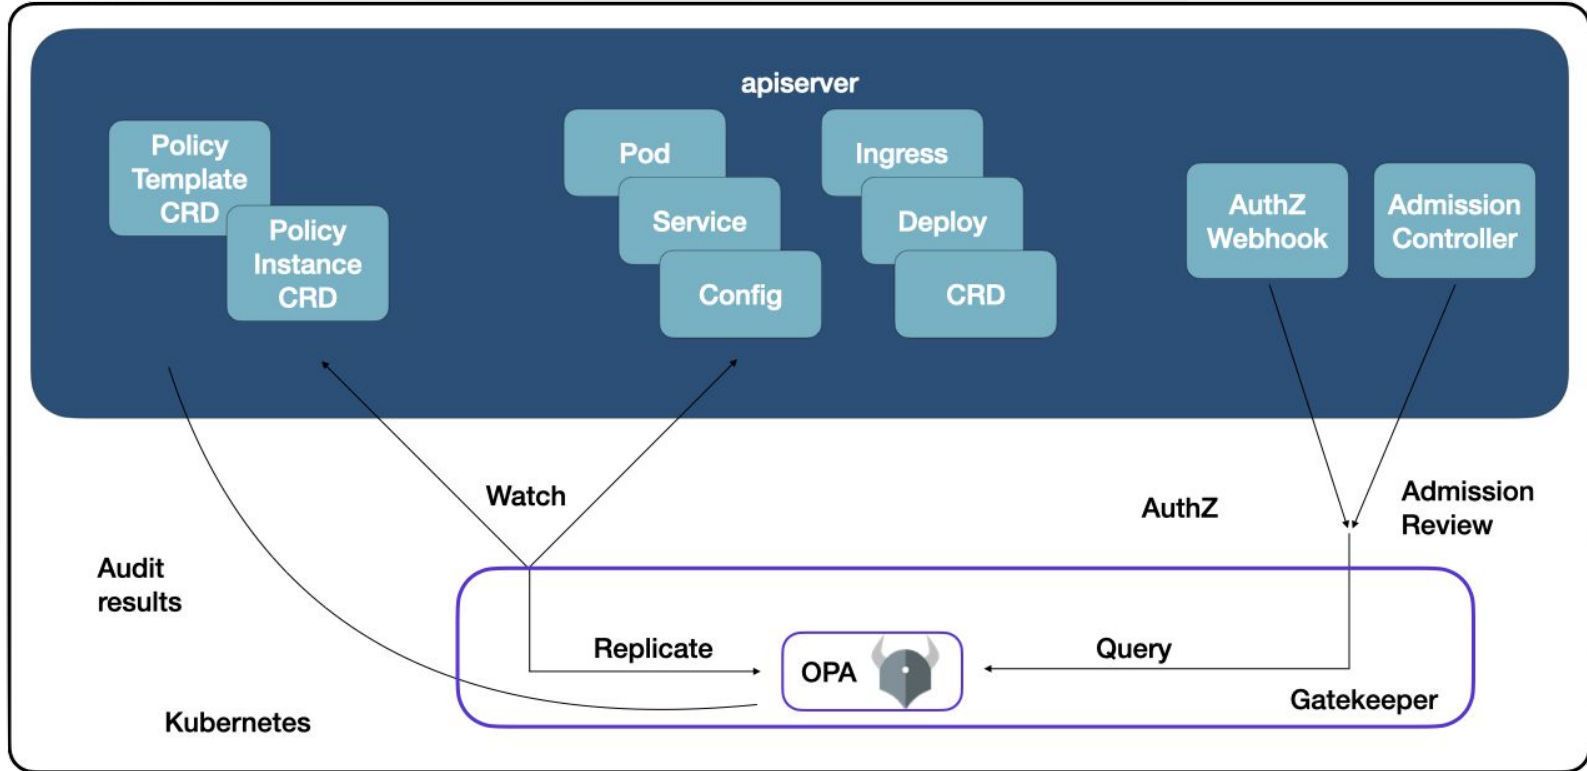
Using **code** to define compliance rules and security policies
is **not very scalable** in the long run.



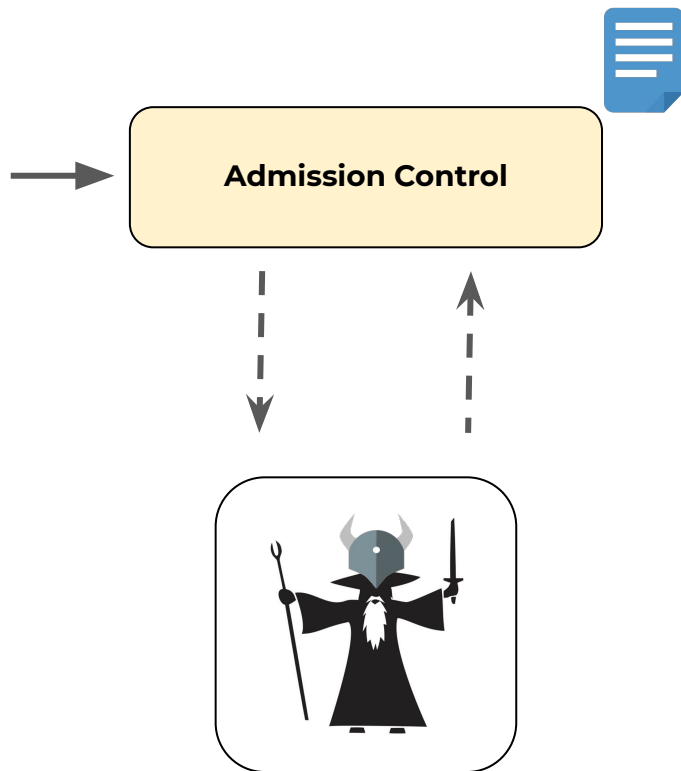
Introducing OPA



OPA Gatekeeper

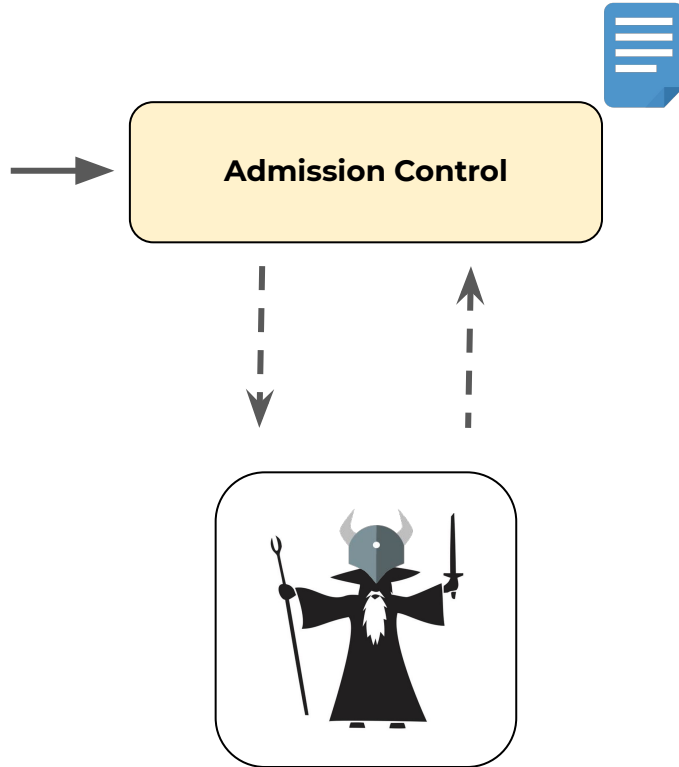


OPA Gatekeeper - WebhookConfiguration



```
apiVersion: admissionregistration.k8s.io/v1
kind: ValidatingWebhookConfiguration
metadata:
  labels:
    gatekeeper.sh/system: "yes"
  name: gatekeeper-validating-webhook-configuration
webhooks:
- admissionReviewVersions:
  - v1
  - v1beta1
  clientConfig:
    service:
      name: gatekeeper-webhook-service
      namespace: gatekeeper-system
      path: /v1/admit
  failurePolicy: Ignore
  matchPolicy: Exact
  name: validation.gatekeeper.sh
  namespaceSelector:
    matchExpressions:
    - key: admission.gatekeeper.sh/ignore
      operator: DoesNotExist
  rules:
  - apiGroups:
    - '*'
    apiVersions:
    - '*'
    operations:
    - CREATE
    - UPDATE
    resources:
    - '*'
  sideEffects: None
  timeoutSeconds: 3
```

OPA Gatekeeper - WebhookConfiguration



```
apiVersion: admissionregistration.k8s.io/v1
kind: ValidatingWebhookConfiguration
metadata:
  labels:
    gatekeeper.sh/system: "yes"
  name: gatekeeper-validating-webhook-configuration
webhooks:
- admissionReviewVersions:
  - v1
  - v1beta1
  clientConfig:
    service:
      name: gatekeeper-webhook-service
      namespace: gatekeeper-system
      path: /v1/admit
  failurePolicy: Ignore
  matchPolicy: Exact
  name: validation.gatekeeper.sh
  namespaceSelector:
    matchExpressions:
    - key: admission.gatekeeper.sh/ignore
      operator: DoesNotExist
  rules:
  - apiGroups:
    - '*'
    apiVersions:
    - '*'
    operations:
    - CREATE
    - UPDATE
    resources:
    - '*'
  sideEffects: None
  timeoutSeconds: 3
```

```
package k8srequiredlabels

deny[{"msg": msg, "details": {"missing_labels": missing}}] {
    provided := {label | input.review.object.metadata.labels[label]}
    required := {label | label := input.parameters.labels[_]}
    missing := required - provided
    count(missing) > 0
    msg := sprintf("you must provide labels: %v", [missing])
}
```

OPA Constraint Framework - ConstraintTemplate



```
apiVersion: templates.gatekeeper.sh/v1beta1
kind: ConstraintTemplate
metadata:
  name: k8srequiredlabels
spec:
  crd:
    spec:
      names:
        kind: K8sRequiredLabels
        listKind: K8sRequiredLabelsList
        plural: k8srequiredlabels
        singular: k8srequiredlabels
      validation:
        # Schema for the `parameters` field
        openAPIV3Schema:
          properties:
            labels:
              type: array
              items: string
  targets:
    - target: admission.k8s.gatekeeper.sh
      rego: |
```

```
package k8srequiredlabels
```

```
deny[{"msg": msg, "details": {"missing_labels": missing}}] {
  provided := {label | input.review.object.metadata.labels[label]}
  required := {label | label := input.parameters.labels[_]}
  missing := required - provided
  count(missing) > 0
  msg := sprintf("you must provide labels: %v", [missing])
}
```

```
apiVersion: constraints.gatekeeper.sh/v1beta1
kind: K8sRequiredLabels
metadata:
  name: ns-must-have-hr
spec:
  match:
    kinds:
      - apiGroups: [""]
        kinds: ["Namespace"]
  parameters:
    labels: ["hr"]
```

OPA Constraint Framework - Audits



```
apiVersion: constraints.gatekeeper.sh/v1beta1
kind: K8sRequiredLabels
metadata:
  name: ns-must-have-hr
spec:
  #...
status:
  auditTimestamp: "2019-08-06T01:46:13Z"
  byPod:
    - enforced: true
      id: gatekeeper-controller-manager-0
  violations:
    - enforcementAction: deny
      kind: Namespace
      message: 'you must provide labels: {"hr"}'
      name: default
    - enforcementAction: deny
      kind: Namespace
      message: 'you must provide labels: {"hr"}'
      name: gatekeeper-system
```


The fall (aka the demo)



1. Deploy OPA Gatekeeper
2. Deploy Gatekeeper Policy Manager
3. Deploy simple policies
4. Audit violations

Thank you





That's it!



@niccoloraspa



/in/niccolo-raspa/



niccolo@sighup.io



Slides and Demo available at **bit.ly/kcd-italy-opa**